

Investigating TMDb dataset

Introduction

TMDb dataset is a dataset that contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue.

The columns selected for the context of this report are: "id", "popularity", "budget", "revenue", "cast", "genres"

Data Wrangling

General Properties

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv('./tmdb-movies.csv')
df.head()
```

```
Out[2]:
```

	id	imdb_id	popularity	budget	revenue	original_title	cast	
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...	http://www.thed
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D...	http://ww

	id	imdb_id	popularity	budget	revenue	original_title	cast
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	Vin Diesel Paul Walker Jason Statham Michelle ...

5 rows × 21 columns



In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     10866 non-null  int64
1   imdb_id                10856 non-null  object
2   popularity              10866 non-null  float64
3   budget                 10866 non-null  int64
4   revenue                 10866 non-null  int64
5   original_title          10866 non-null  object
6   cast                   10790 non-null  object
7   homepage                2936 non-null  object
8   director               10822 non-null  object
9   tagline                 8042 non-null  object
10  keywords                9373 non-null  object
11  overview                10862 non-null  object
12  runtime                 10866 non-null  int64
13  genres                  10843 non-null  object
14  production_companies    9836 non-null  object
15  release_date            10866 non-null  object
16  vote_count              10866 non-null  int64
17  vote_average            10866 non-null  float64
18  release_year            10866 non-null  int64
19  budget_adj              10866 non-null  float64
20  revenue_adj             10866 non-null  float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

In [4]: `df.isnull().sum()`

```
Out[4]: id                     0
imdb_id                    10
popularity                  0
budget                      0
revenue                     0
original_title              0
cast                        76
homepage                   7930
director                    44
tagline                     2824
keywords                   1493
overview                     4
runtime                      0
genres                      23
production_companies       1030
release_date                0
vote_count                  0
```

```
vote_average      0
release_year      0
budget_adj        0
revenue_adj       0
dtype: int64
```

```
In [5]: df.dtypes
```

```
Out[5]: id                int64
imdb_id              object
popularity          float64
budget              int64
revenue             int64
original_title      object
cast                object
homepage            object
director            object
tagline             object
keywords            object
overview            object
runtime             int64
genres              object
production_companies object
release_date        object
vote_count          int64
vote_average        float64
release_year        int64
budget_adj          float64
revenue_adj         float64
dtype: object
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	id	popularity	budget	revenue	runtime	vote_count	vote_aver
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000	10866.000000	10866.000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863	217.389748	5.974
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405	575.619058	0.935
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000	10.000000	1.500
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000	17.000000	5.400
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000	38.000000	6.000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000	145.750000	6.600
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000	9767.000000	9.200

Remove unnecessary columns

```
In [7]: df.drop(['imdb_id', 'homepage', 'director', 'overview', 'tagline', 'keywords',
                'budget_adj', 'revenue_adj', 'production_companies', 'original_title', 'release_d
                axis=1, inplace=True)
df.head()
```

```
Out[7]:
```

	id	popularity	budget	revenue	cast	runtime	genres	vote_c
--	----	------------	--------	---------	------	---------	--------	--------

	id	popularity	budget	revenue	cast	runtime	genres	vote_c
0	135397	32.985763	150000000	1513528810	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	124	Action Adventure Science Fiction Thriller	
1	76341	28.419936	150000000	378436354	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	120	Action Adventure Science Fiction Thriller	
2	262500	13.112507	110000000	295238201	Shailene Woodley Theo James Kate Winslet Ansel...	119	Adventure Science Fiction Thriller	
3	140607	11.173104	200000000	2068178225	Harrison Ford Mark Hamill Carrie Fisher Adam D...	136	Action Adventure Science Fiction Fantasy	
4	168259	9.335014	190000000	1506249360	Vin Diesel Paul Walker Jason Statham Michelle ...	137	Action Crime Thriller	

Remove rows with null values

```
In [8]: df.isna().sum()
```

```
Out[8]: id                0
popularity              0
budget                 0
revenue                0
cast                   76
runtime                0
genres                 23
vote_count             0
vote_average           0
release_year           0
dtype: int64
```

```
In [9]: df.dropna(inplace=True)
df.isna().sum()
```

```
Out[9]: id                0
popularity              0
budget                 0
revenue                0
cast                   0
runtime                0
genres                 0
vote_count             0
vote_average           0
release_year           0
dtype: int64
```

```
In [10]: df.shape
```

Out[10]: (10768, 10)

Remove duplicates

```
In [11]: df.drop_duplicates(inplace=True)
df.shape
```

Out[11]: (10767, 10)

Fixing the genres column

The genres column contains the genres of each movie separated with "|". To make it easier to work with, the genres cell in each row will be transformed to be a list with genres names then it will be spreaded over multiple rows with one genre per row for each movie in another dataframe

Change the genres field to an array:

```
In [12]: df['genres'] = df['genres'].apply(lambda x: x.split('|'))
```

```
In [13]: df.head(1)
```

Out[13]:

	id	popularity	budget	revenue	cast	runtime	genres	vote_count	vote_ave
0	135397	32.985763	150000000	1513528810	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	124	[Action, Adventure, Science Fiction, Thriller]	5562	

```
In [14]: genres_df = pd.DataFrame(df['genres'].tolist(), index=df['id']).stack()
genres_df = genres_df.reset_index([0, 'id'])
genres_df.columns=['id', 'genres']
genres_df.head()
```

Out[14]:

	id	genres
0	135397	Action
1	135397	Adventure
2	135397	Science Fiction
3	135397	Thriller
4	76341	Action

```
In [15]: df_copy = df.copy()
df_copy.drop(['genres'], axis=1, inplace=True)
genres_df = pd.merge(df_copy, genres_df, how='inner', left_on='id', right_on='id')
genres_df.head()
```

Out[15]:

	id	popularity	budget	revenue	cast	runtime	vote_count	vote_average	release
--	----	------------	--------	---------	------	---------	------------	--------------	---------

	id	popularity	budget	revenue	cast	runtime	vote_count	vote_average	release_date
0	135397	32.985763	150000000	1513528810	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	124	5562	6.5	
1	135397	32.985763	150000000	1513528810	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	124	5562	6.5	
2	135397	32.985763	150000000	1513528810	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	124	5562	6.5	
3	135397	32.985763	150000000	1513528810	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	124	5562	6.5	
4	76341	28.419936	150000000	378436354	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	120	6185	7.1	

Fixing the cast column

The same process on the genres will be applied on the cast

```
In [16]: df['cast'] = df['cast'].apply(lambda x: x.split('|'))
df.head(1)
```

	id	popularity	budget	revenue	cast	runtime	genres	vote_count	vote_average
0	135397	32.985763	150000000	1513528810	[Chris Pratt, Bryce Dallas Howard, Irrfan Khan...	124	[Action, Adventure, Science Fiction, Thriller]	5562	6.5

```
In [17]: cast_df = pd.DataFrame(df['cast'].tolist(), index=df['id'].stack())
cast_df = cast_df.reset_index([0, 'id'])
cast_df.columns=['id', 'cast']
cast_df.head()
```

	id	cast
--	----	------

	id	cast
0	135397	Chris Pratt
1	135397	Bryce Dallas Howard
2	135397	Irrfan Khan
3	135397	Vincent D'Onofrio
4	135397	Nick Robinson

```
In [18]: df_copy = df.copy()
df_copy.drop(['cast'], axis=1, inplace=True)
cast_df = pd.merge(df_copy, cast_df, how='inner', left_on='id', right_on='id')
cast_df.head()
```

```
Out[18]:
```

	id	popularity	budget	revenue	runtime	genres	vote_count	vote_average	release_y
0	135397	32.985763	150000000	1513528810	124	[Action, Adventure, Science Fiction, Thriller]	5562	6.5	2
1	135397	32.985763	150000000	1513528810	124	[Action, Adventure, Science Fiction, Thriller]	5562	6.5	2
2	135397	32.985763	150000000	1513528810	124	[Action, Adventure, Science Fiction, Thriller]	5562	6.5	2
3	135397	32.985763	150000000	1513528810	124	[Action, Adventure, Science Fiction, Thriller]	5562	6.5	2
4	135397	32.985763	150000000	1513528810	124	[Action, Adventure, Science Fiction, Thriller]	5562	6.5	2


Each of the cast and genres dataframes are saved separately away from the main dataframe to not influence the result by the duplications in the other columns

EDA

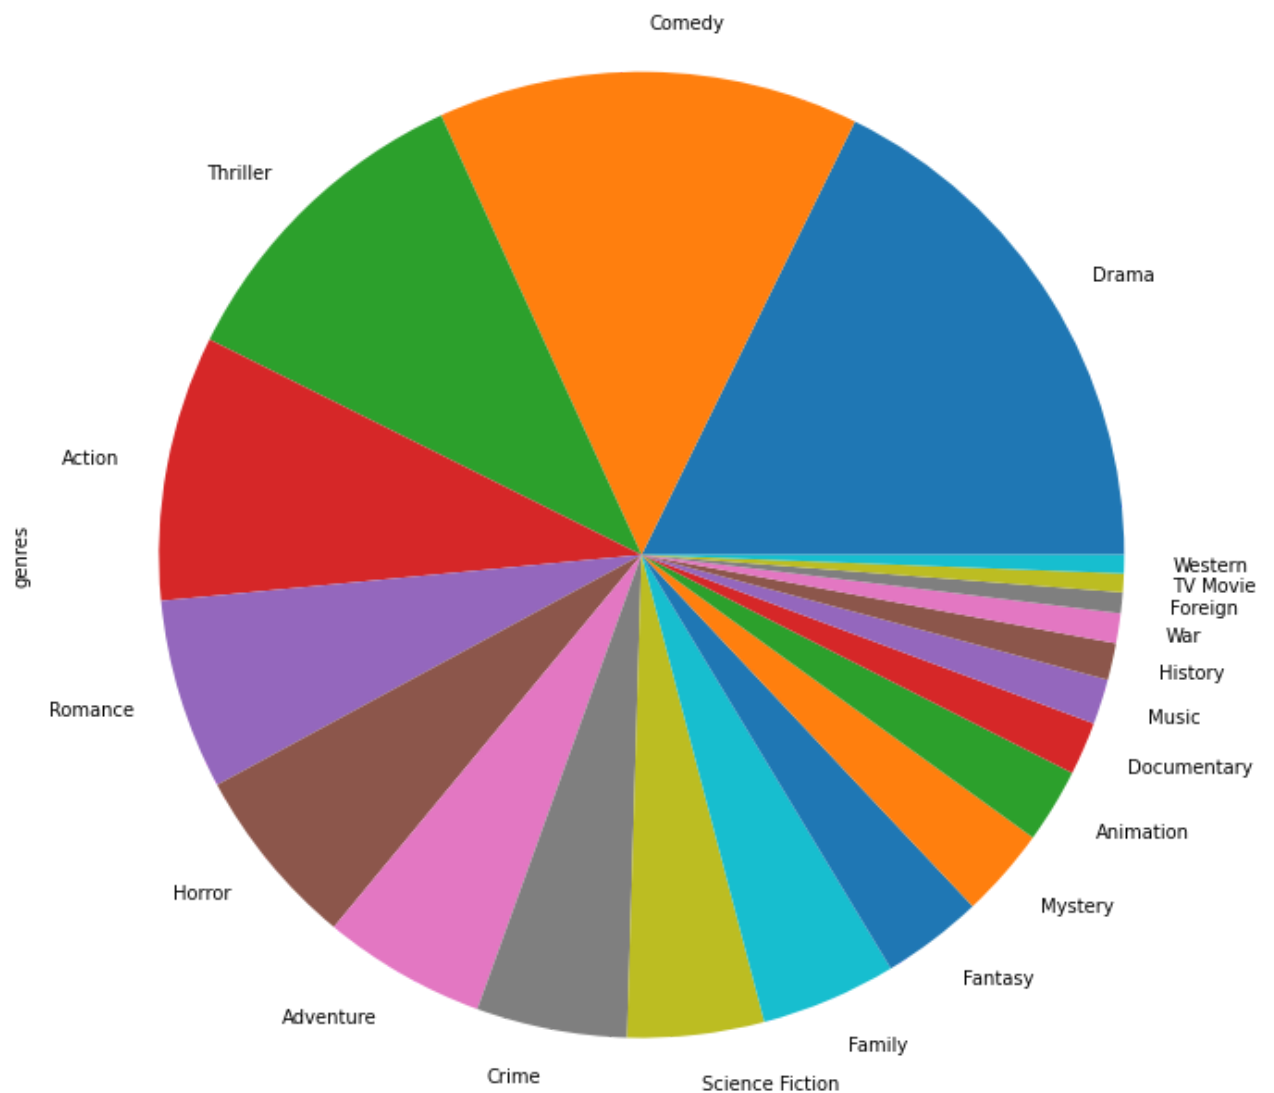
Time to ask questions, find conclusions, and have some fun!

1. What are the most used genres in movies?

```
In [19]: genres_df['genres'].value_counts().plot(kind='bar', figsize=(16,12))  
plt.show()
```



```
In [20]: genres_df['genres'].value_counts().plot(kind='pie', figsize=(16,12))  
plt.show()
```

The most used genres are Drama, Comedy, Thriller, and Action. While TV movie, War, Western, and Foreign, and other genres are lightly produced.

2. What is the average movie budget for each genre?

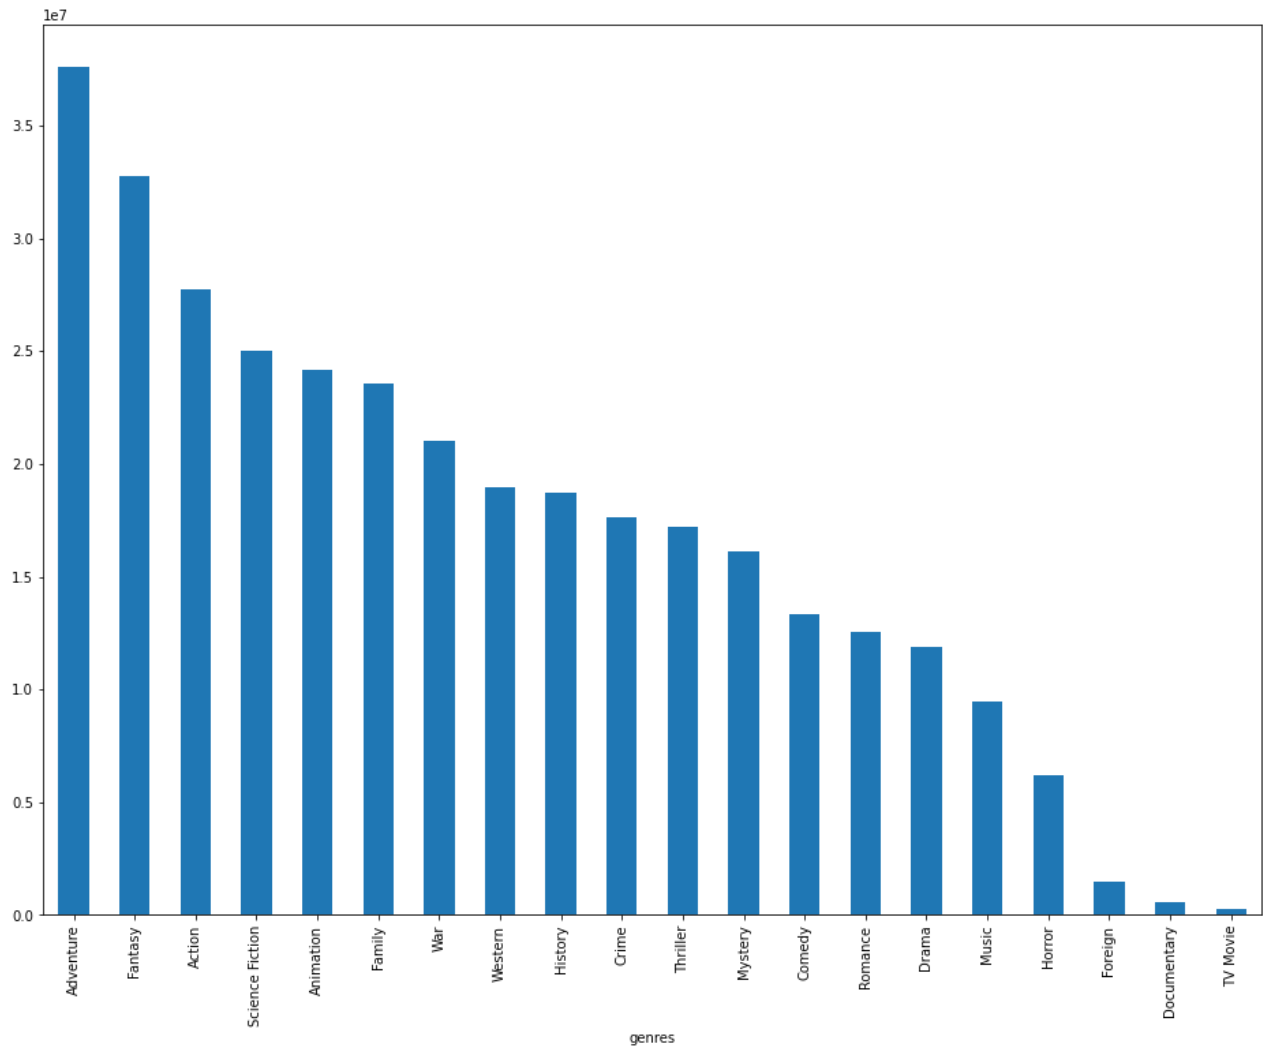
```
In [21]: avg_budget = genres_df.groupby('genres').mean()['budget']

# round the numbers
avg_budget = avg_budget.apply(lambda x: round(x))

# sort the values
avg_budget = avg_budget.sort_values(ascending=False)
avg_budget.head()
```

```
Out[21]: genres
Adventure      37594809
Fantasy        32791579
Action         27762757
Science Fiction 25013386
Animation      24194674
Name: budget, dtype: int64
```

```
In [22]: avg_budget.plot(kind='bar', figsize=(16,12))
plt.show()
```



The most expensive genres are Adventure, Fantasy, Family, History, and Science Fiction. While TV Movie, Foreign, Documentary movies are the least expensive.

3. What is the average movie revenue for each genre?

```
In [23]: avg_revenue = genres_df.groupby('genres').mean()['revenue']

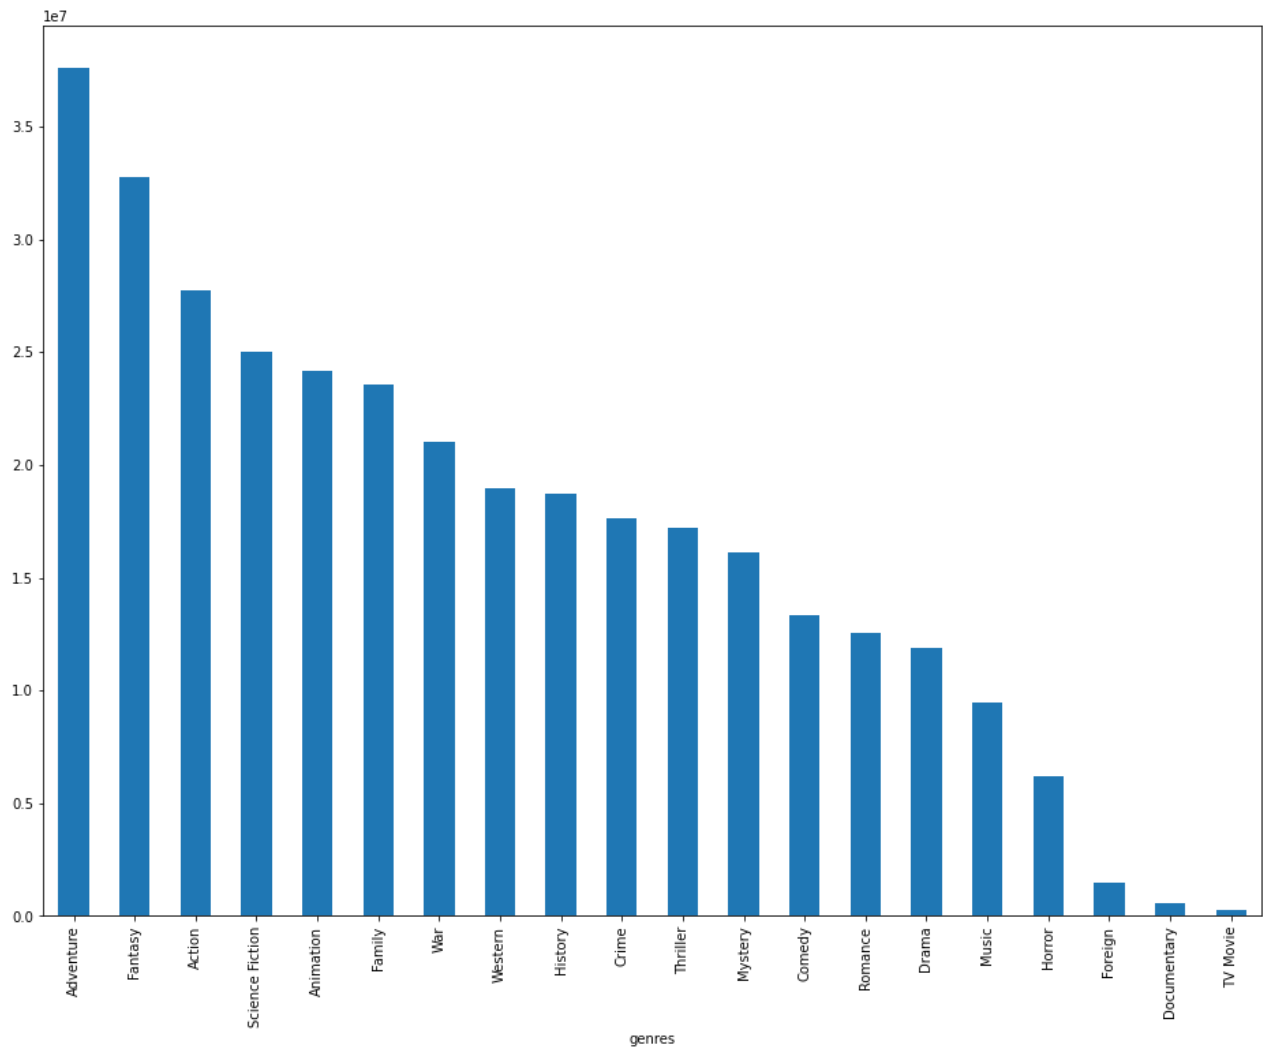
# round the numbers
avg_revenue = avg_revenue.apply(lambda x: round(x))

# sort the values
avg_revenue = avg_revenue.sort_values(ascending=False)
avg_revenue.head()
```

```
Out[23]: genres
Adventure    113291895
Fantasy      96842272
Animation    78630774
Family       73146218
Action       72886452
Name: revenue, dtype: int64
```

```
In [24]: avg_budget.plot(kind='bar', figsize=(16,12))
```

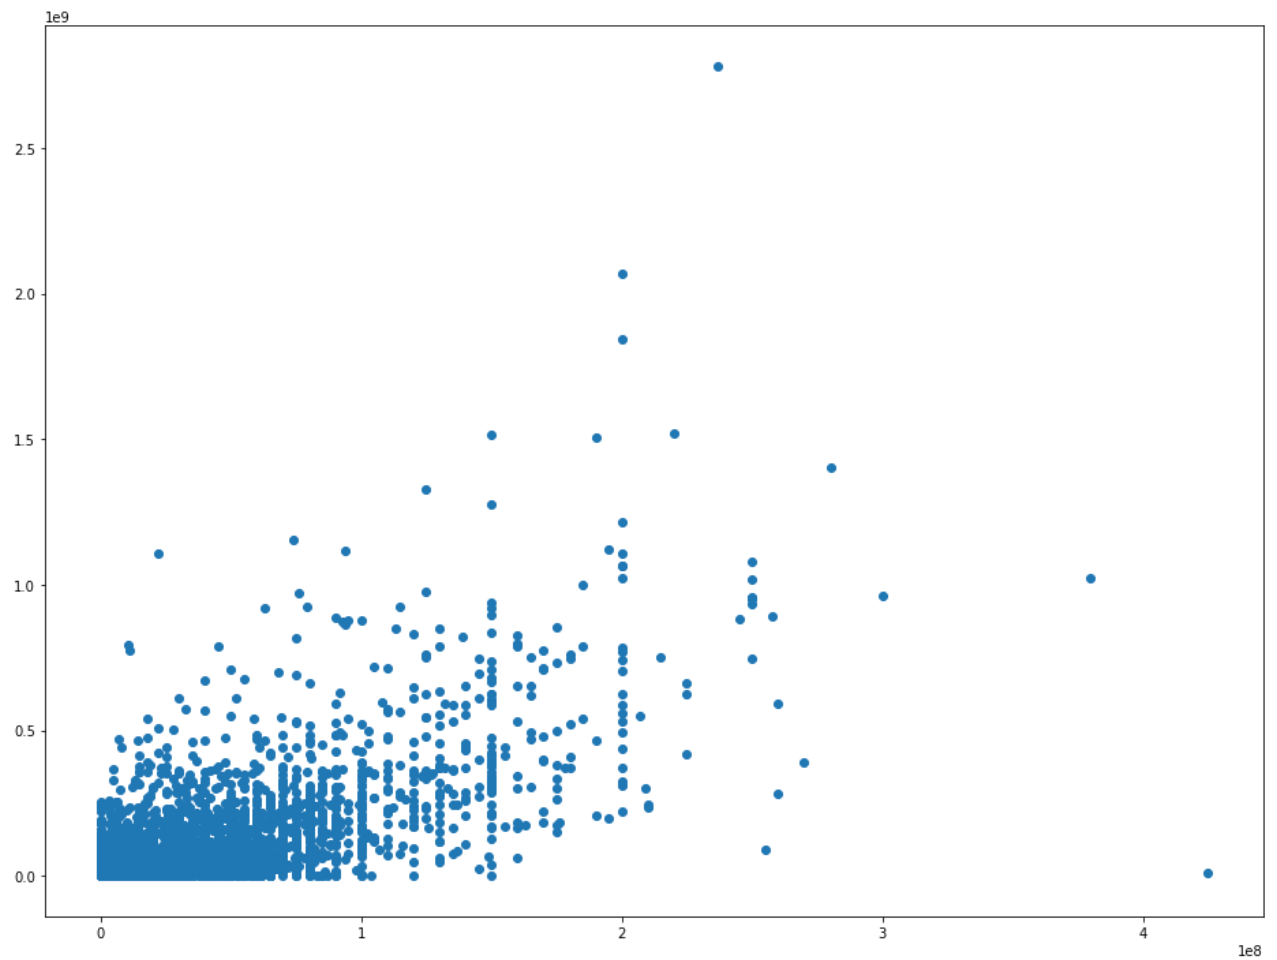
```
plt.show()
```



The genres that generates highest revenues are Adventure, Action, Fantasy, and Family. While Documentary movies don't generate much of a revenue.

4. Does the budget influence the revenue?

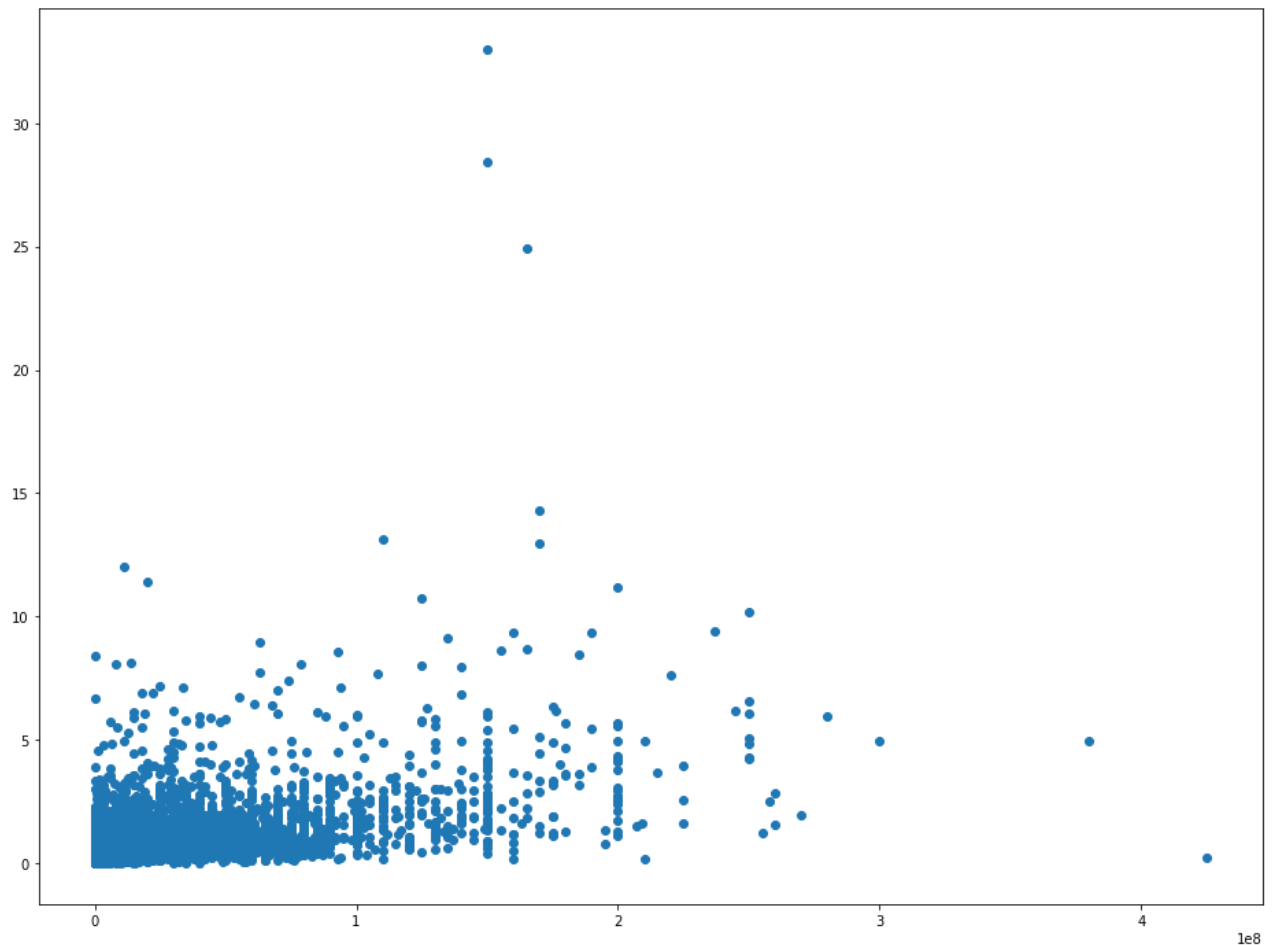
```
In [25]: plt.figure(figsize=(16,12))
plt.scatter(df['budget'], df['revenue'])
plt.show()
```



It appears that there is a very low positive correlation between the budget spent and the revenue gained. In other words, budget doesn't have that large influence on the revenue.

5. Does the budget spent influence the popularity?

```
In [26]: plt.figure(figsize=(16,12))
plt.scatter(df['budget'], df['popularity'])
plt.show()
```



Again, although the budget spent have small influence on the popularity, it still not a large influence to consider.

6. How the cast influence the revenue


```
In [27]: avg_cast_revenue = cast_df.groupby('cast').mean()['revenue']

# round the numbers
avg_cast_revenue = avg_cast_revenue.apply(lambda x: round(x))

# sort the values
avg_cast_revenue = avg_cast_revenue.sort_values(ascending=False)
avg_cast_revenue.head()
```

```
Out[27]: cast
Daisy Ridley          2068178225
Santino Fontana       1274219009
Bob Peck              920100000
Robert MacNaughton    792910554
Rosemary Harris       783766341
Name: revenue, dtype: int64
```

```
In [28]: avg_cast_revenue[:100].plot(kind='bar', figsize=(16,12))
plt.show()
```



Cast members rarely have effect on the revenue of the movies they participate in. Daisy Ridley had the highest effect on the revenue of the movies she participated in while other cast members had similar effect on the movies they participate in.

Conclusions

- The most used genres are Drama, Comedy, Thriller, and Action with the Drama on top. People tend to be "Drama" queens
- The highest budget spent on Adventure, Fantasy, Family, and History.
- The highest revenue making genres are the ones most produces. That didn't need a data analyst to find out.
- Budget spent has low influence on both revenues on popularity. Money doesn't buy everything after all. At least not the popularity of the movies.
- Cast members has no influence on the revenues. It turned out to be just a myth.